

COLOSSEUM · SOLANA

Vitals

A patient is dying on a clock and you decide what happens next.

Anyone can play. Nobody can fake the replay.

EP1 · THE LAST BITE

THE PRODUCT

**A nineteen-year-old is dying on a clock.
You are the doctor. Every run leaves a proof.**

Vitals EP1 · THE LAST BITE · F9DBB7DE05 [← episodes](#) [restart](#) [resume](#) [easy](#) LOCALNET · TREE #40348 · 3 ANCHORED · 0 YOURS



MONITOR Deteriorating

HR
128

BP
83/54

SpO₂
94%

RR
28

[Monitor](#) [Vent](#) [Pump](#)

BED 3 · ER Ing · F 19 · Deteriorating 2 ALARM 14:21:30

SINUS 25 mm/s

PLETH

RESP

HR
128
bpm

SpO₂
94
%

NIBP
88/54
mmHg · 16 s ago

RR
28

TALK TO HER

She answers in her own words, from a model running on our own hardware. She can be wrong, and she can be too breathless to answer.

TREAT HER

Orders, drugs, airway, fluids. A real bedside monitor moves the way the physiology moves.

SHE LIVES OR SHE DOES NOT

No script. A deterministic clinical automaton decides, in real time, from what you did and when.

THE RUN IS PROVABLE

Every case reduces to one hash. Anyone can replay the tape and get the same answer, without asking us.

ALREADY IN PRODUCTION

Eighteen of Thailand's twenty-nine medical faculties already run the engine underneath this.

Vitals is the verifiable layer on **Embla** — a clinical simulation engine live for two months, adopted without a sales team, and national champion at NECTEC AI for Thai 2026. Live right now, in production: embla.megawiz.co.th

18 / 29

Medical faculties with active learners — counting verified faculties, not self-reported signups

290+

Clinicians and students — individual accounts, not institutional seats

671

Scored runs across 433 authored clinical scenarios

1st

Grand Champion, NECTEC AI for Thai 2026

THE TEAM THAT SHIPPED IT

PARIPOL TOOPIROH

CEO & AI architect — Embla core engine, Rust end to end

DR. SIRAWIT TANSIRI

Co-founding physician — clinical accuracy, OSCE rubrics, safety gates

VORANICHAYANAN

Co-founder & COO — registered nurse; hospital ops & SOP co-design

PONGSAKORN JAIKAEW

AI engineer — persona engine, ASR/TTS, automated evals

NAPAPHACH W.

BA & QA lead — scenario quality, E2E, on-site testing

PIMNIPA “MUK”
CHAOSIRIKUL

Final-year medical student, Siriraj — plays every case as the learner it is built for; realism & usability review

The adoption is Embla's, earned before any of this was anchored. What Vitals adds is the part a school cannot do for itself: make the record checkable by someone who does not trust the school.

THE PROBLEM, AND THE MISSION

One doctor for every 1,080 people in Thailand. One for every 2,150 in Indonesia.

Our mission: one per cent more doctors who actually reach patients, every year — *licensed* doctors, not merely graduates.

PEOPLE PER DOCTOR · PHYSICIANS PER 1,000 POPULATION · WHO / WORLD BANK, 2019–20

1:370

Japan · 2.7 per 1,000

1:570

world average · 1.75 per 1,000

1:1,080

Thailand · 0.93 per 1,000

1:2,150

Indonesia · 0.47 per 1,000

A Thai doctor carries 2.9× the load a Japanese one does, and an Indonesian doctor 2.0× the Thai one — the marginal doctor in Jakarta stands in front of nearly six times the people the marginal doctor in Tokyo does. **A product that charges aims itself at the places it matters least.**

WHY 1% IS A TARGET AND NOT A SLOGAN

Attrition is **9.1%** globally — the review that reports it gives a range of **2.7–20.1%**, and a second review puts it at 11.1%. Moving 9.1% to 8.2% yields 1% more graduates, so **1% means cutting attrition by about a tenth** — 1 in 10 of everyone who leaves today, whatever their reason. The reason cited most often is **academic difficulty, at 55.7%**, and it is the one we act on — cited, not apportioned: the four reasons in that review overlap and sum to 145.7%, so no share of dropouts can be read off it. Order-of-magnitude effects already exist here: in one South American cohort (Vergel et al.), curriculum reform cut dropout from **41.5% to 3.3%**. We need a tenth of that.

A number we used to lead with, kept with the wording it needs: across 3,142 *US* counties, every 10 additional *primary care* physicians per 100,000 people was *associated* with 51.5 more days of life expectancy — roughly 1,400 person-years across that whole population (Basu et al., *JAMA Internal Medicine* 2019). It is a US primary-care association measured over populations, not an effect any one doctor produces.

GAMEPLAY & VERIFIABLE OUTPUT

The replay is the trophy, not the score.



GHOST · MOCKUP

you · adrenaline 1:12 ghost · 0:38

Ghost racing. The tape replays exactly, so another player's run can run beside yours — and you watch the two patients diverge while you are still deciding. Trackmania's ghost car, in an emergency room.



RECORD · MOCKUP

harm · stand/walk collapse 50 saves · 12 harms

Harm is on the record. Fifty saves with twelve harms reads differently from thirty with none — and it keeps the fastest run and the cleanest run from being the same run.



DEATH CARD · MOCKUP

DeathArrest · turned at 2:14

gold path → adrenaline IM at 0:40

Failure is content. "She died because you reached for antihistamines" is a better clip than any win. A loss should be legible, teach in the same breath, and be worth showing someone.

treated adrenaline IM · O ₂ · supine · fluids · admit	WinDischarge · leaf: 65c56e...
stood up adrenaline IM · stood the patient up · O ₂ · fluids · admit	WinDischarge · harm · leaf: a6bec3...
no adrenaline antihistamine · steroid · wait	DeathArrest · fatal

Deterministic execution: Three players, three tapes, replayed against the shipped physiology automaton. Same outcome as the first, different leaf — because harm is on the record.

ONE SEASON · EP1 IS FILMED, THE REST ARE AUTHORED AUTOMATA

EP1

The Last Bite

anaphylaxis · she has minutes

EP2

Time Is Muscle

STEMI · the clock is muscle

EP3

Don't Make Him Cry

epiglottitis · doing nothing is the treatment

EP4

The Masquerader


PE · everything looks normal

EP5

The Night the Stars Fell

mass casualty · the queue is the cost

The model makes it playable. The automaton makes it provable. Two inputs in, one hash out — and the program refuses its own user.



The tape and the scenario are the only inputs. The chain never sees a run — only a leaf that has to prove against a root it built itself.

LIVE VALIDATOR · CLAIM: PROFICIENT

REJECTED — claimed Proficient, computed Competent (3 distinct attempts)

LIVE VALIDATOR · CLAIM: COMPETENT

GRANTED — stored Competent, xp 35. Stored level is what program computed.

No issuer key. No oracle. The Solana program recomputes the level from Merkle proofs and writes only what its own arithmetic agrees with.

Free to the learner. Paid by whoever relies on the proof.

The learner never pays — we tested charging, **eighteen faculties use it and none bought**, and charging would aim this at the countries that need it least. The economics live on the other side: **a doctor’s proof that travels** — checkable in any country, by anyone, without asking any institution — and **every act of reliance is a paid, on-chain event**.

THE LOOP — FOUR ACTORS, ONE RAIL

learner	plays, practises, anchors a record	— free —
relying party	residency programme verifies a candidate	pays USDC
case author	publishes a station anyone can train on	earns per attempt
verifier	would replay the tape — bonded, slashed if it lies	earns the fee

More learners → more records worth relying on → more paid verification → more authors earning → more cases → more learners. **The flywheel runs on reliance, not speculation.** Today only the learner row runs: the program holds no money and moves none, and the three paid rows are **designed, not built**.

WHAT ONLY A CHAIN DOES HERE

BORDERLESS BY CONSTRUCTION

The relying party is in another country. A shared neutral layer is the only database both sides can trust — take the chain away and you are back to one company’s word.

DONATIONS THE PROGRAM COULD ENFORCE

Coming next: money given for a scenario could only be spent on that scenario, released against the delivered content hash, so a funder recomputes our impact instead of reading our report. The program holds no money today; that instruction is not written yet.

ROYALTIES THAT CROSS BORDERS

Designed, not built: case authors would earn cents per attempt, settled in seconds — amounts card rails would consume entirely. This is what makes an open case marketplace possible at all, and it is why the Case Registry is a drawing rather than a shipped instruction.

ON THE TOKEN QUESTION, BOTH LAYERS IN ONE BREATH

There is no token. Progression and credentials are **deliberately untradeable** — progression is a record the program owns rather than a mint, so there is nothing to sell, and a market in “Expert in Cardiology” would void the product. That half is built. The other half is a drawing: **\$VIGIL** *would* be a separate verification-market token that **players never hold and never see**, bonding the operators who replay tapes and slashed when one lies, with the arena named **3R** so the two never share a name. **Designed, not built** — no mint, no bond account and no slashing instruction exists in this repository, and any fungible token is an explicit scope cut for this sprint.

MARKET AND PRECEDENT

This shape already won this hackathon once.

FRONTIER 2026 · GRAND CHAMPION


CrowdBrain: *"Train in simulation → Qualify through QA → Route to real work."*

The exact same shape — proven at the largest crypto hackathon ever run (2,857 submissions, 150 countries) for robot operators instead of clinicians.

PATTERN: VERIFIABLE HUMAN IN THE LOOP

HOW EPISODE SIX GETS MADE — SKALD CONSOLE

AI SHOWRUNNER + QA



Skald verifies the content. Solana verifies the play.

Both projects reached the same conclusion independently: **the honest report has to be collateralised**. Their robot nodes are bonded. Ours would be — bonded, slashed when they lie — and that bond is **designed, not built**. What is built is the property that would make it enforceable: replay is exact, so two honest verifiers cannot disagree and a lie is caught by arithmetic.

WHY MEDICINE FIRST

Highest stakes & regulation. Solve trust here and it generalises downward.

WHY SOLANA

~\$110 anchors a million compressed attempts · author royalties would be cents, settled in seconds · the learner is gasless and never sees a seed phrase. The numbers only close on this chain.

HEAD START (2 MOS LIVE)

18 of 29 Thai medical faculties · 290+ clinicians and students · 1st Place @ NECTEC 2026.

WHERE IT STANDS, AND WHAT'S NEXT

Two halves built. Five weeks to join them, then the arena.

WORKING NOW (LIVE IN PRODUCTION)

- **Replay** → **leaf**. EP1 runs against shipped automaton; three tapes, three distinct leaves, reproducible across runs.
- **Program** → **refusal**. Native Solana program **live on public devnet** (535FMHHZ...BTSg), recomputing levels and rejecting invalid claims — deployed bytecode verified byte-for-byte against the build.
- **Award-Winning Engine**. 1st Place @ NECTEC AI for Thai 2026 (team01.aiforthai.in.th/embla.megawiz.co.th).
- **Rust end to end**. One scoring implementation compiled into game, verifier, and onchain program.

Cut on purpose

STATED BEFORE ANYONE ASKS

40 of 100 points are re-derivable; 60 are attested — signers name the model and version, and more signers means more confidence, never proof. **We cannot prove the human at the keyboard was the enrolled student**; identity binding is delegated to the issuer, and the credential carries the issuer's name so a relying party can price that issuer's rigour. **We do not grant standing** — only a board can make competence authoritative. And **the pass-rate number does not exist yet**: what exists is the measurement reliability, the study design, and the date it lands.

ZK selective disclosure · verifier DePIN with live staking · token launch · mobile · new clinical content. Each is a real v2 line and none survives the run-up. The credential layer stays a stub until an accrediting institution is standing next to it — progression doesn't need one, which is why progression is the demo.

THE FIVE WEEKS BEFORE THE ARENA OPENS

COMMIT-REVEAL, IN CODE

The record gains the commitment and the slot it was made at, so the proof travels in the leaf rather than depending on which program version ran.

THE OSCE SCORER, PORTED CLEAN

The deterministic 40 as a dependency-free crate, bound to the reference engine by its own conformance vectors. The judged 60 stays outside the auditable core, where it belongs.

THE UNLOCK THE CHAIN CHECKS

Stars earned → the program recomputes the predicate → the next scenario opens. Legible without knowing any medicine.

LIVE, PUBLIC, AND METERED

vitals.academy open to strangers with the first scenario free and no account — because four weeks of real usage cannot be written the week before a deadline.

The 60-second version: you play a patient who is dying, the chain keeps a replay of what you did, and anyone can re-derive it without asking us.

The ask: the people who would actually rely on a record like this. We can build the protocol — we cannot grant it standing.